

# **Vectorized VByte Decoding**

Jeff Plaisance, Nathan Kurz, Daniel Lemire

# **Inverted Indexes**

# Inverted Index

- Like index in the back of a book
- words = terms, page numbers = doc ids
- Term list is sorted
- Doc list for each term is sorted

# Standard Index

doc id	query	country	impressions	clicks
0	software	Canada	10	1
1	blank	Canada	10	0
2	sales	US	5	0
3	software	US	8	1
4	blank	US	10	1

# Constructing an Inverted Index

	query			country		impression			clicks	
doc id	blank	sales	software	Canada	US	5	8	10	0	1
0			✓	✓				✓		✓
1	✓			✓				✓	✓	
2		✓			✓	✓			✓	
3			✓		✓		✓			✓
4	✓				✓			✓		✓

# Constructing an Inverted Index

field	term	0	1	2	3	4
query	blank		✓			✓
	sales			✓		
	software	✓			✓	
country	Canada	✓	✓			
	US			✓	✓	✓
impressions	5			✓		
	8				✓	
	10	✓	✓			✓
clicks	0		✓	✓		
	1	✓			✓	✓

# Inverted Index

field	term	doc list
query	blank	1, 4
	sales	2
	software	0, 3
country	Canada	0, 1
	US	2, 3, 4
impressions	5	2
	8	3
	10	0, 1, 4
clicks	0	1, 2
	1	0, 3, 4

# Inverted Indexes

Allow you to:

- Quickly find all documents containing a term
- Intersect several terms to perform boolean queries



# Inverted Index Optimizations

- Compressed data structures
  - Better use of RAM and processor cache
  - Better use of memory bandwidth
  - Increased CPU usage and time
- Optimizations matter!

# Delta / VByte Encoding

- Doc id lists are sorted
- Delta between a doc id and the previous doc id is sufficient
- Deltas are usually small integers

# Delta Encoding

field	term	doc list
query	nursing	34, 86, 247, 301, 674, 714

# Delta Encoding

field	term	doc list
query	nursing	34, 86, 247, 301, 674, 714
		34, 52, 161, 54, 373, 40

# Small Integer Compression

- Golomb/Rice
- VByte (or Varint)
- Binary Packing
- PForDelta

# Small Integer Compression

- Golomb/Rice
- **VByte**
- Bit Packing
- PForDelta

# VByte Encoding

9838

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---



# VByte Encoding

9838

0	0	0	0	0	0	0	0
0	0	1	0	0	1	1	0

0	0	0	0	0	0	0	0
0	1	1	0	1	1	1	0

# VByte Encoding

9838

0	0	0	0	0	0	0	0
0	0	1	0	0	1	1	0

0	0	0	0	0	0	0	0
0	1	1	0	1	1	1	0

# VByte Encoding

9838

0	0	0	0	0	0	0	0
0	0	1	0	0	1	1	0

0	0	0	0	0	0	0	0
0	1	1	0	1	1	1	0

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

?	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

?	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

?	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

1	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

1	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---



# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

1	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

?	1	0	0	1	1	0	0
---	---	---	---	---	---	---	---

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

1	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

?	1	0	0	1	1	0	0
---	---	---	---	---	---	---	---

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

1	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

?	1	0	0	1	1	0	0
---	---	---	---	---	---	---	---

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

1	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

0	1	0	0	1	1	0	0
---	---	---	---	---	---	---	---

# VByte Encoding

9838

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---

0	0	1	0	0	1	1	0
---	---	---	---	---	---	---	---

0	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

1	1	1	0	1	1	1	0
---	---	---	---	---	---	---	---

0	1	0	0	1	1	0	0
---	---	---	---	---	---	---	---

# VByte

Pros:

- Compression
- Can fit more of index in RAM
- Higher information throughput per byte read from disk

# VByte

Cons:

- Decodes one byte at a time
- Lots of branch mispredictions
- Not fast to decode
- Largest ints expand to 5 bytes

# Vectorized VByte Decoding

Optimized decoder implemented using x86\_64  
intrinsics



# Vectorized VByte Decoding

```
01001010 11001000 01110001 01001110
10011011 01101010 10110101 00010111
01110110 10001101 10110011 11000001
```

# Vectorized VByte Decoding

01001010	11001000	01110001	01001110
10011011	01101010	10110101	00010111
01110110	10001101	10110011	11000001

pmovmskb: Extract top bit of each byte

# Vectorized VByte Decoding

01001010	11001000	01110001	01001110
10011011	01101010	10110101	00010111
01110110	10001101	10110011	11000001

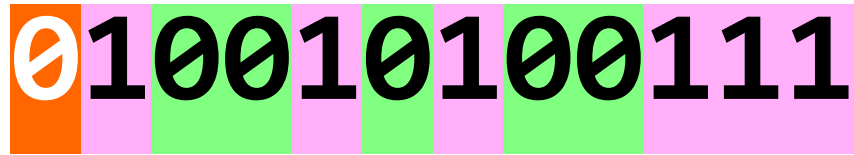
pmovmskb: Extract top bit of each byte

010010100111

**010010100111**

Pattern of leading bits determines:

- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume



010010100111

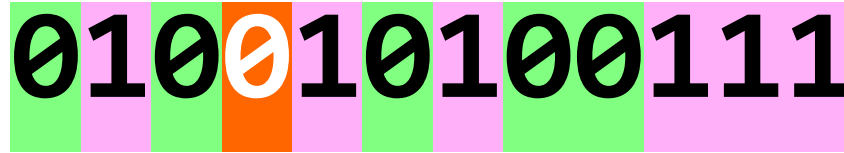
Pattern of leading bits determines:

- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume

010010100111

Pattern of leading bits determines:

- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume



010010100111

Pattern of leading bits determines:

- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume

010010100111

Pattern of leading bits determines:

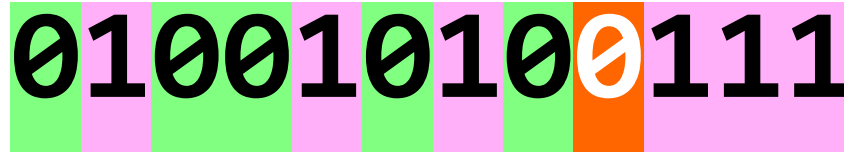
- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume



010010100111

Pattern of leading bits determines:

- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume



010010100111

Pattern of leading bits determines:

- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume

010010100111

Pattern of leading bits determines:

- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume

010010100111

Pattern of leading bits determines:

- how many varints to decode
- sizes and offsets of varints
- length of longest varint in bytes
- number of bytes to consume

010010100111

Decoding options for:

- sixteen 1 byte varints
- six 1-2 byte varints
- four 1-3 byte varints
- two 1-5 byte varints

010010100111

Decoding options for:

- sixteen 1 byte varints - special case
- six 1-2 byte varints -  $2^6$ , 64 possibilities
- four 1-3 byte varints -  $3^4$ , 81 possibilities
- two 1-5 byte varints -  $5^2$ , 25 possibilities

170 total possibilities

010010100111

## Data Distribution:

- Longer doc id lists are necessarily composed of smaller deltas
- Most deltas in real datasets (ClueWeb09, Indeed's internal datasets) fall into 1 byte case or 1-2 byte case

# Most Significant Bit Decoding

- We separate most significant bit decoding from integer decoding
- Reduces duplicate most significant bit decoding work if we don't consume all 12 bytes
- Better instruction level parallelism



010010100111

- If most significant bits of next 16 bytes are all 0, handle sixteen 1 byte ints case
- Otherwise lookup most significant bits of next 12 bytes in 4096 entry lookup table

010010100111

Lookup table contains:

- Shuffle vector index from 0-169 representing which possibility we are decoding
- Number of bytes of input that will be consumed

010010100111

Branch on shuffle vector index to determine which case we are decoding

- 0-63 - six 1-2 byte ints
- 64-144 - four 1-3 byte ints
- 145-169 - two 1-5 byte ints

# Six 1-2 Byte Ints

01001010	11001000	01110001	01001110
10011011	01101010	10110101	00010111
01110110	10001101	10110011	11000001

Decode 6 varints from 9 bytes

# Expected Positions

Six 1-2 byte ints

2	1	2	1	2	1	2	1	2	1	2	1	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Four 1-3 byte ints

0	3	2	1	0	3	2	1	0	3	2	1	0	3	2	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Two 1-5 byte ints

1	5	0	4	0	3	0	2	1	5	0	4	0	3	0	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

# Six 1-2 Byte Ints

01001010	11001000	01110001	01001110
10011011	01101010	10110101	00010111
01110110	10001101	10110011	11000001

Pad out 1 byte ints to 2 bytes

# Six 1-2 Byte Ints

01001010 00000000

01001110 00000000

10110101 00010111

11001000 01110001

10011011 01101010

01110110 00000000

Pad out 1 byte ints to 2 bytes

# Shuffle input

- Use index to lookup appropriate shuffle vector
- Shuffle input bytes to get them in the expected positions



# pshufb

```
for (i = 0; i < 16; i++) {  
    if (mask[i] & 0x80) {  
        dest[i] = 0;  
    } else {  
        dest[i] = src[mask[i] & 0xF];  
    }  
}
```

# pshufb

src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

# pshufb

src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

# pshufb

src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

[illegible]



# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

## dest

[illegible]

# src

mask

dest[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]



# pshufb

src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]



# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

[illegible]

# pshufb

# src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

# dest

[illegible]



# pshufb

src

DE	DF	27	E3	7C	A9	60	55	1C	EA	45	56	A6	43	C9	48
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

mask

0	11	2	-1	12	-1	13	4	0	10	2	6	4	3	13	5
---	----	---	----	----	----	----	---	---	----	---	---	---	---	----	---

dest

DE	56	27	0	A6	0	43	7C	DE	45	27	60	7C	E3	43	A9
----	----	----	---	----	---	----	----	----	----	----	----	----	----	----	----

# Shuffle input

- Use index to lookup appropriate shuffle vector
- Shuffle input bytes to get them in the expected positions

# Six 1-2 Byte Ints

01001010	00000000	11001000	01110001
01001110	00000000	10011011	01101010
10110101	00010111	01110110	00000000

Reverse bytes in 2 byte varints

\*not actually necessary since x86 is little endian



# Six 1-2 Byte Ints

00000000	01001010	01110001	11001000
00000000	01001110	01101010	10011011
00010111	10110101	00000000	01110110

Reverse bytes in 2 byte varints

\*not actually necessary since x86 is little endian

# Six 1-2 Byte Ints

00000000	01001010	01110001	11001000
00000000	01001110	01101010	10011011
00010111	10110101	00000000	01110110

Mask out leading purple 1's

## Six 1-2 Byte Ints

00000000	01001010	01110001	01001000
00000000	01001110	01101010	00011011
00010111	00110101	00000000	01110110

Mask out leading purple 1's

# Six 1-2 Byte Ints

00000000	01001010	01110001	01001000
00000000	01001110	01101010	00011011
00010111	00110101	00000000	01110110

Shift top bytes of each varint 1 bit right  
(mask/shift/or)

## Six 1-2 Byte Ints

00000000 01001010	00111000 11001000
00000000 01001110	00110101 00011011
00001011 10110101	00000000 01110110

Shift top bytes of each varint 1 bit right  
(mask/shift/or)

# Six 1-2 Byte Ints

00000000	01001010	00111000	11001000
00000000	01001110	00110101	00011011
00001011	10110101	00000000	01110110

Done!

# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

101101000110



# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

101101000110

# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

101101000110

# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

101101000110

# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

101101000110

# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

101101000110

# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

Decode 4 varints from 8 bytes

# Four 1-3 Byte Ints

11101110	00011101	11110101	11101101
01111001	11111000	01101001	00100001
00001011	10110101	10111001	01110110

Pad ints to 4 bytes

# Four 1-3 Byte Ints

11101110	00011101	00000000	00000000
11110101	11101101	01111001	00000000
11111000	01101001	00000000	00000000
00100001	00000000	00000000	00000000

Pad ints to 4 bytes



# Four 1-3 Byte Ints

00000000	00000000	00011101	11101110
00000000	01111001	11101101	11110101
00000000	00000000	01101001	11111000
00000000	00000000	00000000	00100001

Reverse bytes

\*not actually necessary since x86 is little endian

# Four 1-3 Byte Ints

00000000	00000000	00011101	11101110
00000000	01111001	11101101	11110101
00000000	00000000	01101001	11111000
00000000	00000000	00000000	00100001

Clear top bit of each byte

# Four 1-3 Byte Ints

00000000	00000000	00011101	01101110
00000000	01111001	01101101	01110101
00000000	00000000	01101001	01111000
00000000	00000000	00000000	00100001

Clear top bit of each byte

# Four 1-3 Byte Ints

00000000	00000000	00011101	01101110
00000000	01111001	01101101	01110101
00000000	00000000	01101001	01111000
00000000	00000000	00000000	00100001

Shift 2nd least significant bytes over by 1 bit

(mask/shift/or)

# Four 1-3 Byte Ints

00000000	00000000	00001110	11101110
00000000	01111001	00110110	11110101
00000000	00000000	00110100	11111000
00000000	00000000	00000000	00100001

Shift 2nd least significant bytes over by 1 bit

(mask/shift/or)

# Four 1-3 Byte Ints

00000000	00000000	00001110	11101110
00000000	01111001	00110110	11110101
00000000	00000000	00110100	11111000
00000000	00000000	00000000	00100001

Shift 3rd least significant bytes over by 2 bits

(mask/shift/or)

# Four 1-3 Byte Ints

00000000	00000000	00001110	11101110
00000000	00011110	01110110	11110101
00000000	00000000	00110100	11111000
00000000	00000000	00000000	00100001

Shift 3rd least significant bytes over by 2 bits

(mask/shift/or)

# Four 1-3 Byte Ints

00000000	00000000	00001110	11101110
00000000	00011110	01110110	11110101
00000000	00000000	00110100	11111000
00000000	00000000	00000000	00100001

Done!



# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

111101110110

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

111101110110

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

111101110110

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

111101110110

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

Decode 2 varints from 9 bytes

## Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

- Could handle the same way as other cases
- Would require 5 AND operations, 4 shift operations, and 4 OR operations

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

- We can simulate shifting by different amounts with multiplication
- Only needs 1 multiplication, 1 shift, 1 OR, 1 shuffle

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

Two 1-5 byte ints

1	5	0	4	0	3	0	2	1	5	0	4	0	3	0	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Treat SIMD register as eight 16 bit registers, loading 1 byte into each. First byte doesn't need to be shifted.



# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

00000000	00000000	00000000	00000000
00000000	00000000	00000000	00000000

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

11101110	00000000	00000000	00000000
00000000	00000000	00000000	00000000

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110
11101110	00000000	00000000	00000000
00000000	00000000	00000000	10011101

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110
11101110	00000000	00000000	00000000
00000000	11110101	00000000	10011101

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

11101110	00000000	00000000	11101101
00000000	11110101	00000000	10011101

# Two 1-5 Byte Ints

11101110	10011101	11110101	11101101
00000011	11111000	11101001	10100001
00001011	10110101	10111001	01110110

11101110	00000011	00000000	11101101
00000000	11110101	00000000	10011101

# Two 1-5 Byte Ints

11101110	00000011
00000000	11101101
00000000	11110101
00000000	10011101

Clear top bit of each byte

# Two 1-5 Byte Ints

01101110	00000011
00000000	01101101
00000000	01110101
00000000	00011101

Clear top bit of each byte



## Two 1-5 Byte Ints

01101110 00000011 \* 16 (<< 4)

00000000 01101101 \* 32 (<< 5)

00000000 01110101 \* 64 (<< 6)

00000000 00011101 \* 128 (<< 7)

Multiply to shift bits into place

# Two 1-5 Byte Ints

01101110	00000011	*	16	( << 4 )
00000000	01101101	*	32	( << 5 )
00000000	01110101	*	64	( << 6 )
00000000	00011101	*	128	( << 7 )

Multiply to shift bits into place

# Two 1-5 Byte Ints

11100000 00110000 \* 16 (<< 4)

00001101 10100000 \* 32 (<< 5)

00011101 01000000 \* 64 (<< 6)

00001110 10000000 \* 128 (<< 7)

Multiply to shift bits into place

# Two 1-5 Byte Ints

11100000 00110000

00011101 01000000

00001101 10100000

00001110 10000000

# Two 1-5 Byte Ints

1110	0000	0011	0000	0000	1101	101	00000
000	11101	01	0000000	00	001110	1	00000000

Left shift everything by 8 bits

# Two 1-5 Byte Ints

1110	0000	0011	0000	0000	1101	101	00000
000	11101	01	0000000	00	001110	1	00000000

Left shift everything by 8 bits

0011	0000	0000	1101	101	00000	000	11101
01	0000000	00	001110	1	00000000	00000000	

# Two 1-5 Byte Ints

1110	0000	0011	0000	0000	1101	101	00000
000	11101	01	0000000	00	001110	1	00000000

Bitwise OR pre-shifted and shifted registers

0011	0000	0000	1101	101	00000	000	11101
01	0000000	00	001110	1	00000000	00000000	

# Two 1-5 Byte Ints

1110	0000	0011	0000	0000	1101	101	00000
000	11101	01	0000000	00	001110	1	00000000

Bitwise OR pre-shifted and shifted registers

0011	0000	0000	1101	101	00000	000	11101
01	0000000	00	001110	1	00000000	00000000	



# Two 1-5 Byte Ints

1111	0000	0011	0000	0000	1101	101	00000
000	11101	01	0000000	00	001110	1	00000000

Bitwise OR pre-shifted and shifted registers

0011	0000	0000	1101	101	00000	000	11101
01	0000000	00	001110	1	00000000	00000000	

# Two 1-5 Byte Ints

11110000	00110000	00001101	10100000
00011101	01000000	00001110	10000000

Bitwise OR pre-shifted and shifted registers

00110000	00001101	10100000	00011101
01000000	00001110	10000000	00000000

# Two 1-5 Byte Ints

11110000 00111101 10101101 10100000  
00011101 01000000 00001110 10000000

Bitwise OR pre-shifted and shifted registers

00110000 00001101 10100000 00011101  
01000000 00001110 10000000 00000000

# Two 1-5 Byte Ints

11110000 00111101 10101101 10100000  
00011101 01000000 00001110 10000000

Bitwise OR pre-shifted and shifted registers

00110000 00001101 10100000 00011101  
01000000 00001110 10000000 00000000

# Two 1-5 Byte Ints

11110000 00111101 10101101 10111101  
01011101 01000000 00001110 10000000

Bitwise OR pre-shifted and shifted registers

00110000 00001101 10100000 00011101  
01000000 00001110 10000000 00000000

# Two 1-5 Byte Ints

11110000 00111101 10101101 10111101  
01011101 01000000 00001110 10000000

Bitwise OR pre-shifted and shifted registers

00110000 00001101 10100000 00011101  
01000000 00001110 10000000 00000000

# Two 1-5 Byte Ints

11110000 00111101 10101101 10111101  
01011101 01001110 10001110 10000000

Bitwise OR pre-shifted and shifted registers

00110000 00001101 10100000 00011101  
01000000 00001110 10000000 00000000

# Two 1-5 Byte Ints

11110000 00111101 10101101 10111101  
01011101 01001110 10001110 10000000

Extract result from every other byte



# Two 1-5 Byte Ints

11110000 00111101 10101101 10111101  
01011101 01001110 10001110 10000000

Extract result from every other byte

# Two 1-5 Byte Ints

11110000 00111101 10101101 10111101  
01011101 01001110 10001110 10000000

Extract result from every other byte

# Two 1-5 Byte Ints

00110000 00111101 10101101 10111101  
01011101 01001110 10001110 10000000

Extract result from every other byte

# Two 1-5 Byte Ints

00110000 00111101 10101101 10111101  
01011101 01001110 10001110 10000000

Extract result from every other byte

# Two 1-5 Byte Ints

00110000 00111101 10101101 10111101  
01011101 01001110 10001110 10000000

Extract result from every other byte

# Two 1-5 Byte Ints

00110000 00111101 10101101 10111101  
01011101 01001110 10001110 10000000

Extract result from every other byte

# Two 1-5 Byte Ints

00110000	00111101	10101101	10111101
01011101	01001110	10001110	10000000

Extract result from every other byte

# Two 1-5 Byte Ints

00110000	00111101	10101101	10111101
01011101	01001110	10001110	10000000

Extract result from every other byte



## Two 1-5 Byte Ints

00110000	00111101	10101101	10111101
01011101	01001110	10001110	10000000

Extract result from every other byte

## Two 1-5 Byte Ints

00110000	00111101	10101101	10111101
01011101	01001110	10001110	10000000

OR in low 7 bits of least significant byte

(remember that we stored it in most significant byte position originally)

## Two 1-5 Byte Ints

00110000	00111101	10101101	10111101
01011101	01001110	10001110	11101110

OR in low 7 bits of least significant byte

(remember that we stored it in most significant byte position originally)

## Two 1-5 Byte Ints

00111101 10111101 01001110 11101110

Final result!

## Two 1-5 Byte Ints

00111101 10111101 01001110 11101110

Final result!

Checking my work against initial varint:

11101110 10011101 11110101 11101101

00000011

## Two 1-5 Byte Ints

00111101 10111101 01001110 11101110

Final result!

Checking my work against initial varint:

11101110 10011101 11110101 11101101  
00000011

## Two 1-5 Byte Ints

00111101 10111101 01001110 11101110

Final result!

Checking my work against initial varint:

11101110 10011101 11110101 11101101  
00000011

## Two 1-5 Byte Ints

00111101 10111101 01001110 11101110

Final result!

Checking my work against initial varint:

11101110 10011101 11110101 11101101  
00000011



## Two 1-5 Byte Ints

00111101 10111101    01001110 11101110

Final result!

Checking my work against initial varint:

11101110 10011101    11110101 11101101

00000011

## Two 1-5 Byte Ints

00111101 10111101 01001110 11101110

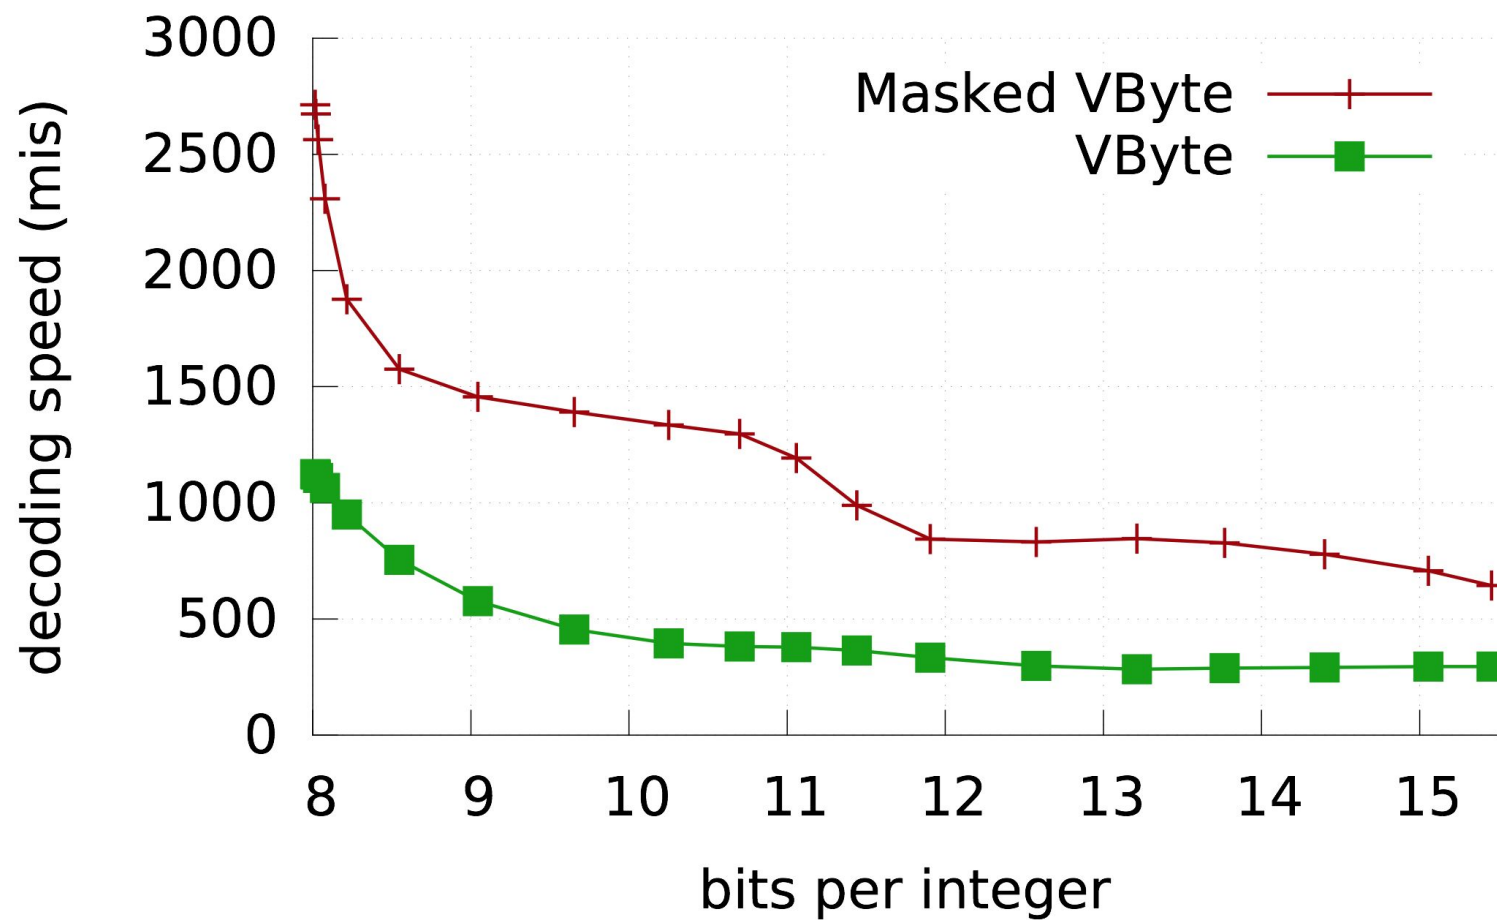
Final result!

Checking my work against initial varint:

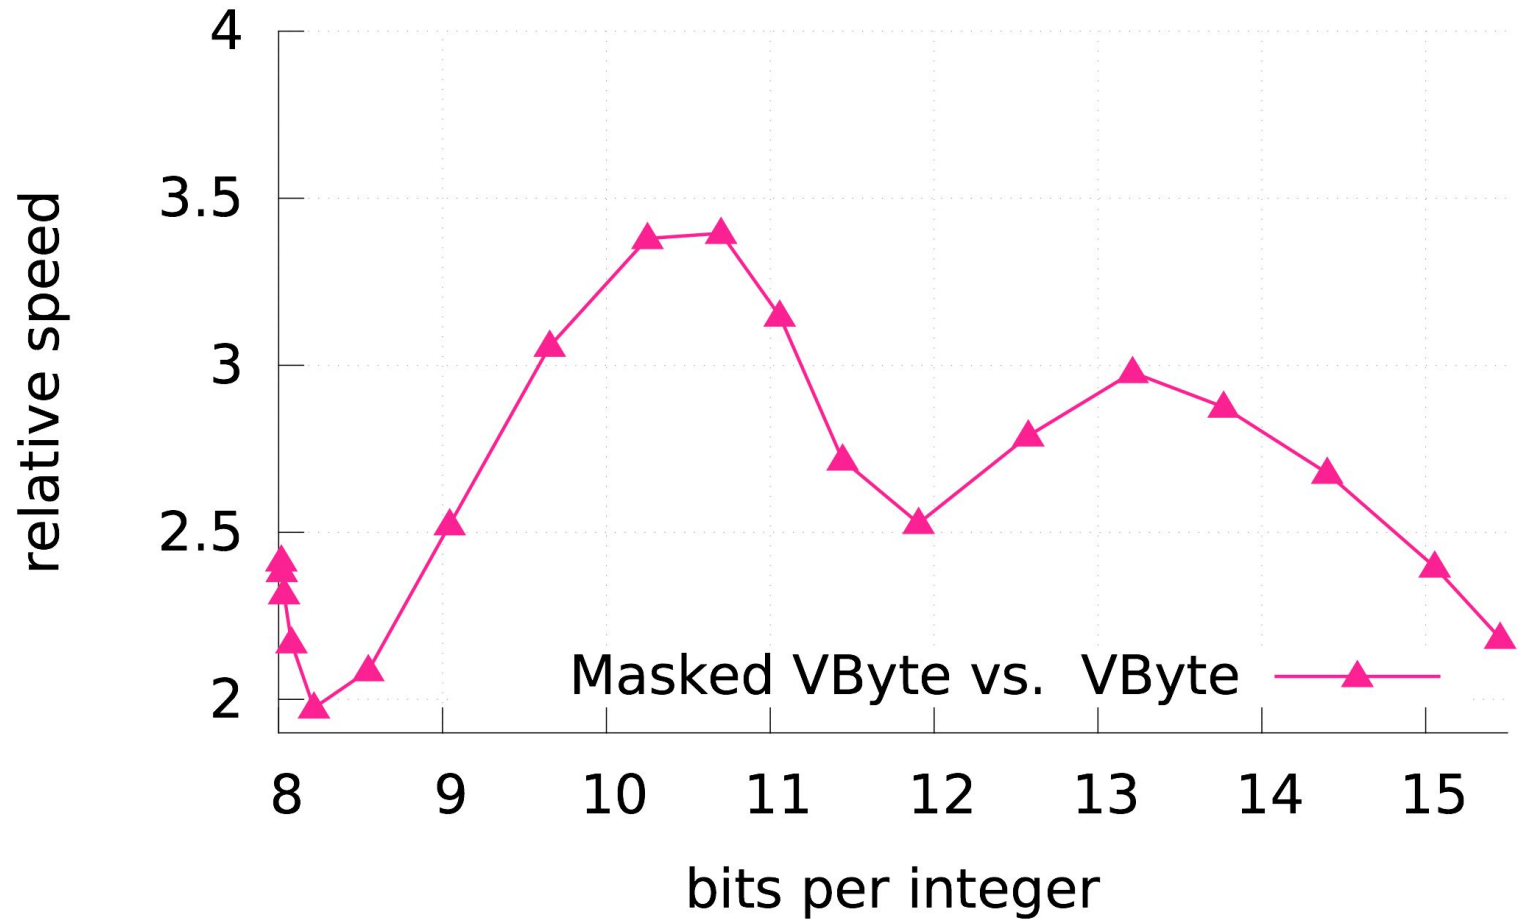
11101110 10011101 11110101 11101101

00000011

# Results



# Results



**Q&A**