

OPEN DATA SCIENCE CONFERENCE



@ODSC

Boston | May 1 - 4 2018

Next Generation Indexes For Big Data Engineering

Daniel Lemire and collaborators

blog: <https://lemire.me>

twitter: [@lemire](https://twitter.com/lemire)

Université du Québec (TÉLUQ)

Montreal 🇨🇦



“One Size Fits All”: An Idea Whose Time Has Come and Gone (Stonebraker, 2005)

Rediscover Unix

In 2018, Big Data Engineering is made of several specialized and re-usable components:

- Calcite : SQL + optimization
- Hadoop
- etc.

"Make your own database engine from parts"

We are in a Cambrian explosion, with thousands of organizations and companies building their custom high-speed systems.

- Specialized used cases
- Heterogeneous data (not everything is in your Oracle DB)

For high-speed in data engineering you need...

- High-level optimizations (e.g., Calcite)
- Indexes (e.g., Pilosa, Elasticsearch)
 - Great compression routines
 - Specialized data structures
-

Sets

A fundamental concept (sets of documents, identifiers, tuples...)

→ For performance, we often work with sets of integers (identifiers).

→ Often 32-bit integers suffice (local identifiers).

- tests : $x \in S$?
- intersections : $S_2 \cap S_1$, unions : $S_2 \cup S_1$, differences : $S_2 \setminus S_1$
- Similarity (Jaccard/Tanimoto): $|S_1 \cap S_1| / |S_1 \cup S_2|$
- Iteration


```
for x in S do  
    print(x)
```


How to implement sets?

- sorted arrays (`std::vector<uint32_t>`)
- hash tables (`java.util.HashSet<Integer>` , `std::unordered_set<uint32_t>`)
- ...
- bitmap (`java.util.BitSet`)
- ❤️ ❤️ ❤️ compressed bitmaps ❤️ ❤️ ❤️

Arrays are your friends

```
while (low <= high) {  
    int mI =  
        (low + high) >>> 1;  
    int m = array.get(mI);  
    if (m < key) {  
        low = mI + 1;  
    } else if (m > key) {  
        high = mI - 1;  
    } else {  
        return mI;  
    }  
}  
return -(low + 1);
```



Hash tables

- value x at index $h(x)$
- random access to a value in expected constant-time
 - much faster than arrays

in-order access is kind of terrible

- [15, 3, 0, 6, 11, 4, 5, 9, 12, 13, 8, 2, **1**, 14, 10, 7]
- [15, 3, 0, 6, 11, 4, 5, 9, 12, 13, 8, **2**, 1, 14, 10, 7]
- [15, **3**, 0, 6, 11, 4, 5, 9, 12, 13, 8, 2, 1, 14, 10, 7]
- [15, 3, 0, 6, 11, **4**, 5, 9, 12, 13, 8, 2, 1, 14, 10, 7]
- [15, 3, 0, 6, 11, 4, **5**, 9, 12, 13, 8, 2, 1, 14, 10, 7]
- [15, 3, 0, **6**, 11, 4, 5, 9, 12, 13, 8, 2, 1, 14, 10, 7]

(Robin Hood, linear probing, MurmurHash3 hash function)

Set operations on hash tables

```
h1 <- hash set
h2 <- hash set
...
for(x in h1) {
  insert x in h2 // cache miss?
}
```

"Crash" Swift

```
var S1 = Set<Int>(1...size)
var S2 = Set<Int>()
for i in d {
    S2.insert(i)
}
```




Swift

Some numbers: half an hour for 64M keys

size	time (s)
1M	0.8
8M	22
64M	1400

- Maps and sets can have quadratic-time performance
<https://lemire.me/blog/2017/01/30/maps-and-sets-can-have-quadratic-time-performance/>
- Rust hash iteration+reinsertion
<https://accidentallyquadratic.tumblr.com/post/153545455987/rust-hash-iteration-reinsertion>



Bitmaps

Efficient way to represent sets of integers.

For example, 0, 1, 3, 4 becomes `0b11011` or "27".

- $\{0\} \rightarrow 0b00001$
- $\{0, 3\} \rightarrow 0b01001$
- $\{0, 3, 4\} \rightarrow 0b11001$
- $\{0, 1, 3, 4\} \rightarrow 0b11011$

Manipulate a bitmap

64-bit processor.

Given x , word index is $x/64$ and bit index $x \% 64$.

```
add(x) {  
    array[x / 64] |= (1 << (x % 64))  
}
```

How fast is it?

```
index = x / 64      -> a shift  
mask = 1 << ( x % 64) -> a shift  
array[ index ] |- mask -> a OR with memory
```

One bit every ≈ 1.65 cycles because of superscalarity

Bit parallelism

Intersection between $\{0, 1, 3\}$ and $\{1, 3\}$
a single AND operation
between `0b1011` and `0b1010` .

Result is `0b1010` or $\{1, 3\}$.

No branching!

Bitmaps love ❤️ wide registers

- SIMD: Single Instruction Multiple Data
 - SSE (Pentium 4), ARM NEON 128 bits
 - AVX/AVX2 (256 bits)
 - AVX-512 (512 bits)

AVX-512 is now available (e.g., from Dell!) with Skylake-X processors.

Bitsets can take too much memory

{1, 32000, 64000} : 1000 bytes for three values

We use compression!

Git (GitHub) utilise EWAH

Run-length encoding

Example: 00000000111111100 est
00000000 — 11111111 — 00

Code long runs of 0s or 1s efficiently.

<https://github.com/git/git/blob/master/ewah/bitmap.c>

Complexity

- Intersection : $O(|S_1| + |S_2|)$ or $O(\min(|S_1|, |S_2|))$
- In-place union ($S_2 \leftarrow S_1 \cup S_2$): $O(|S_1| + |S_2|)$ or $O(|S_2|)$

Roaring Bitmaps

<http://roaringbitmap.org/>

- Apache Lucene, Solr et Elasticsearch, Metamarkets' Druid, Apache Spark, Apache Hive, Apache Tez, Netflix Atlas, LinkedIn Pinot, InfluxDB, Pilosa, Microsoft Visual Studio Team Services (VSTS), Couchbase's Bleve, Intel's Optimized Analytics Package (OAP), Apache Hivemall, eBay's Apache Kylin.
- Java, C, Go (interoperable)

Hybrid model

Set of containers

- sorted arrays ({1,20,144})
- bitset (0b10000101011)
- runs ([0,10],[15,20])

Related to: O'Neil's RIDBit + BitMagic-

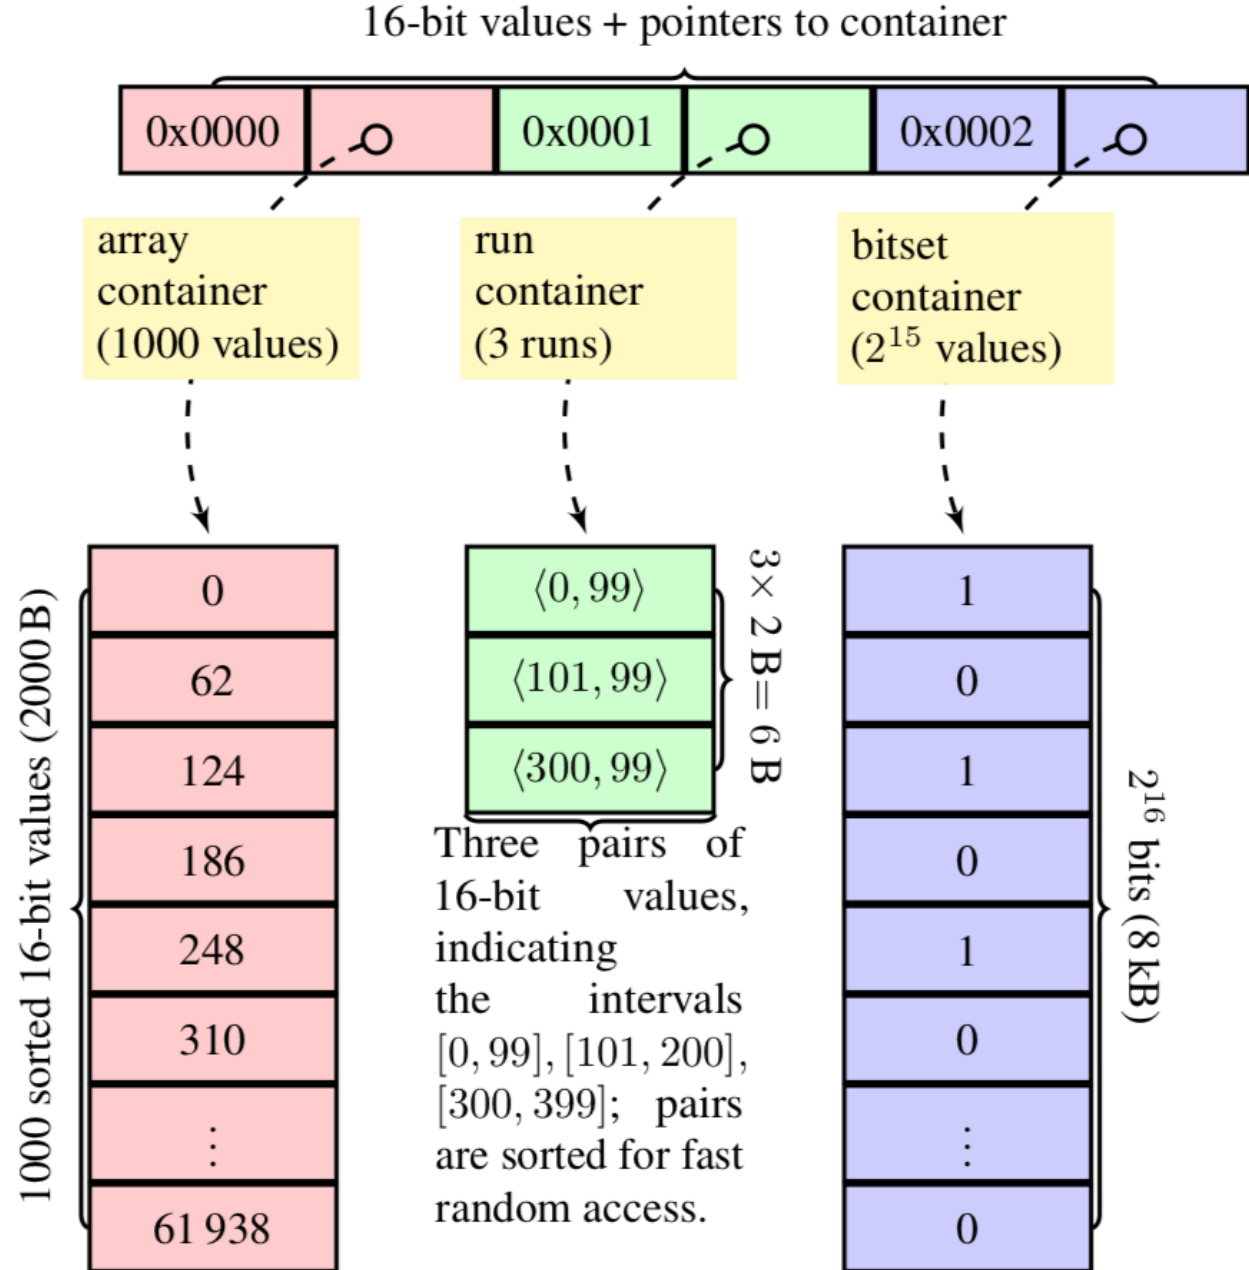


Figure 1. Roaring bitmap containing the first 1000 multiples of 62, all integers in the intervals $[2^{16}, 2^{16} + 100)$, $[2^{16} + 101, 2^{16} + 201)$, $[2^{16} + 300, 2^{16} + 400)$ and all even integers in $[2 \times 2^{16}, 3 \times 2^{16})$.

Roaring

- All containers are small (8 kB), fit in CPU cache
- We predict the output container type during computations
- E.g., when array gets too large, we switch to a bitset
- Union of two large arrays is materialized as a bitset...
- Dozens of heuristics... sorting networks and so on

Use Roaring for bitmap compression whenever possible. Do not use other bitmap compression methods (Wang et al., SIGMOD 2017)

Unions of 200 bitmaps (cycles per input input value)

	bitset	array	hash table	Roaring
census1881	524	32	195	15.1
weather	15.3	32	195	5.38

	bitset	array	hash table	Roaring
census1881	9.85	542	1010	2.6
weather	0.35	94	237	0.16

Integer compression

- "Standard" technique: VByte, VarInt, VInt
- Use 1, 2, 3, 4, ... byte per integer
- Use one bit per byte to indicate the length of the integers in bytes
- Lucene, Protocol Buffers, etc.

varint-GB from Google

- VByte: one branch per integer
- varint-GB: one branch per 4 integers
- each 4-integer block is preceded by a *control byte*

Vectorisation

- Stepanov (STL in C++) working for Amazon proposed varint-G8IU
- Use vectorization (SIMD)
- *Patented*
- Fastest byte-oriented compression technique (until recently)

SIMD-Based Decoding of Posting Lists, CIKM 2011

https://stepanovpapers.com/SIMD_Decoding_TR.pdf

Observations from Stepanov et al.

- We can vectorize Google's varint-GB, but it is not as fast as varint-G8IU


Stream VByte

- Reuse varint-GB from Google
- But instead of mixing control bytes and data bytes, ...
- We store control bytes separately and consecutively...

Daniel Lemire, Nathan Kurz, Christoph Rupp

Stream VByte: Faster Byte-Oriented Integer Compression

Information Processing Letters 130, 2018



Stream VByte is used by...

- Redis (within RediSearch) <https://redislabs.com>
- upscaledb <https://upscaledb.com>
- Trinity <https://github.com/phaistos-networks/Trinity>

Dictionary coding

Use, e.g., by Apache Arrow

Given a list of values:

- "Montreal", "Toronto", "Boston", "Montreal", "Boston"...

Map to integers

- 0, 1, 2, 0, 2

Compress integers:

- Given 2^n distinct values...
- Can use n -bit per values (binary packing, patched coding, frame-of-reference)

Dictionary coding + SIMD

dict. size	bits per value	scalar	AVX2 (256-bit)	AVX-512 (512-bit)
32	5	8	3	1.5
1024	10	8	3.5	2
65536	16	12	5.5	4.5

(cycles per value decoded)

<https://github.com/lemire/dictionary>

To learn more...

- Blog (twice a week) : <https://lemire.me/blog/>
- GitHub: <https://github.com/lemire>
- Home page : <https://lemire.me/en/>
- CRSNG : *Faster Compressed Indexes On Next-Generation Hardware* (2017-2022)
- Twitter  @lemire

